
Cloud & Automation Class

Minikube II

Adam Lahbib - Skander Soltane - Mouna Rouini - Mariem
Cherif

3/27/2023 @ INSAT

Contents

1	Lab Introduction	3
1.1	Objectives	3
1.2	Architecture	3
1.3	Requirements	4
2	Lab Walkthrough	5
2.0.1	Starting Minikube	5
2.0.2	Enabling Addons	5
2.0.3	Cluster Ready	6
2.1	Creating NGINX	6
2.1.1	Creating a service and statefulset	8
2.1.2	Adding HTML to pods	8
2.1.2.1	Curling the service internally	8
2.1.2.2	Creating a simple hello world homepage	9
2.1.2.3	Curling the service again	9
2.2	Persistence and Disaster Recovery	10
3	Conclusion	12

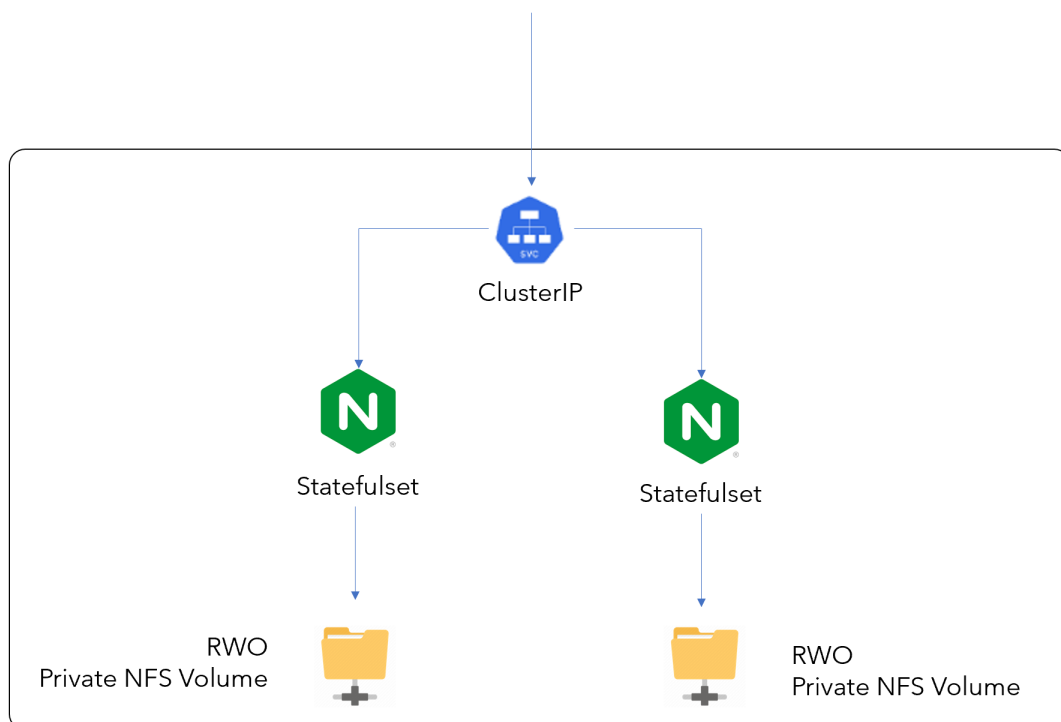
1 Lab Introduction

1.1 Objectives

This lab shows a step-by-step guide to create a web front-end in Kubernetes using a statefulset and a service. Designed to help a developer create a work environment from scratch. This lab emphasizes the features and characteristics of Kubernetes as an orchestrator for container-based workloads.

1.2 Architecture

- A service is created to expose the pods to the outside world.
- A statefulset is created to create two pods connecting to the volume claim template.
- HTML files are stored in a volume claim template, two ReadWriteOnce private NFS volumes.



1.3 Requirements

- Ubuntu 20.04
- Minikube
- Kubectl
- Docker CLI

2 Lab Walkthrough

2.0.1 Starting Minikube

```
adam@4n6nk8s:~$ minikube start --vm-driver='docker'
minikube v1.29.0 on Ubuntu 20.04
✳ Using the docker driver based on user configuration
✳ Using Docker driver with root privileges
🔥 Starting control plane node minikube in cluster minikube
📡 Pulling base image ...
📡 Downloading Kubernetes v1.26.1 preload ...
> preloaded-images-k8s-v18-v1...: 397.05 MiB / 397.05 MiB 100.00% 196.04      > gcr.io/k8s-minikube/kicbase...: 407.19 MiB / 407.19 MiB 100.00%
73.87 M
🔥 Creating docker container (CPUs=2, Memory=2200MB) ...
📡 Preparing Kubernetes v1.26.1 on Docker 20.10.23 ...
  ✳ Generating certificates and keys ...
  ✳ Booting up control plane ...
  ✳ Configuring RBAC rules ...
🔗 Configuring bridge CNI (Container Networking Interface) ...
  ✳ Using image gcr.io/k8s-minikube/storage-provisioner:v5
📡 Verifying Kubernetes components...
🌟 Enabled addons: storage-provisioner, default-storageclass
🔥 Done! Kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

Minikube with the Docker driver is the recommended way to run Kubernetes locally. The Docker driver is the default driver for Minikube. It uses Docker to run Kubernetes components inside a VM on your local machine.

2.0.2 Enabling Addons

```
adam@4n6nk8s:~$ minikube addons enable metrics-server
🔥 metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
  ✳ Using image registry.k8s.io/metrics-server/metrics-server:v0.6.2
🌟 The 'metrics-server' addon is enabled
adam@4n6nk8s:~$ minikube addons enable ingress
🔥 ingress is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
  ✳ Using image registry.k8s.io/ingress-nginx/controller:v1.5.1
  ✳ Using image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
  ✳ Using image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
📡 Verifying ingress addon...
🌟 The 'ingress' addon is enabled
```

```
adam@4n6nk8s:~$ minikube addons list
```

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	3rd party (Ambassador)
auto-pause	minikube	disabled	Google
cloud-spanner	minikube	disabled	Google
csi-hostpath-driver	minikube	disabled	Kubernetes
dashboard	minikube	disabled	Kubernetes
default-storageclass	minikube	enabled	Kubernetes
efk	minikube	disabled	3rd party (Elastic)
freshpod	minikube	disabled	Google
gcp-auth	minikube	disabled	Google
gvisor	minikube	disabled	Google
headlamp	minikube	disabled	3rd party (kinvolk.io)
helm-tiller	minikube	disabled	3rd party (Helm)
inaccel	minikube	disabled	3rd party (InAccel [info@inaccel.com])
ingress	minikube	enabled	Kubernetes
ingress-dns	minikube	disabled	Google
istio	minikube	disabled	3rd party (Istio)
istio-provisioner	minikube	disabled	3rd party (Istio)
kong	minikube	disabled	3rd party (Kong HQ)
kubevirt	minikube	disabled	3rd party (KubeVirt)
logviewer	minikube	disabled	3rd party (unknown)
metallb	minikube	disabled	3rd party (MetalLB)
metrics-server	minikube	enabled	Kubernetes
nvidia-driver-installer	minikube	disabled	Google
nvidia-gpu-device-plugin	minikube	disabled	3rd party (Nvidia)
olm	minikube	disabled	3rd party (Operator Framework)
pod-security-policy	minikube	disabled	3rd party (unknown)
portainer	minikube	disabled	3rd party (Portainer.io)
registry	minikube	disabled	Google
registry-aliases	minikube	disabled	3rd party (unknown)
storage-provisioner	minikube	enabled	Google
storage-provisioner-gluster	minikube	disabled	3rd party (Gluster)
volumesnapshots	minikube	disabled	Kubernetes

```
adam@4n6nk8s:~$
```

The metrics-server is a cluster-wide aggregator of resource usage data. Metrics Server collects resource metrics from Kubelets and exposes them in Kubernetes apiserver through the metrics API for use by Horizontal Pod Autoscaler and Vertical Pod Autoscaler. Metrics API can also be accessed by `kubectl top`, allowing Kubernetes users to check the amount of resources used by their applications.

The ingress addon is a collection of resources that allow you to expose HTTP and HTTPS routes from outside the cluster to services within the cluster. Traffic routing is controlled by rules defined on the Ingress resource.

2.0.3 Cluster Ready

```
adam@4n6nk8s:~$ kubectl get no
```

NAME	STATUS	ROLES	AGE	VERSION
minikube	Ready	control-plane	5m34s	v1.26.1

```
adam@4n6nk8s:~$
```

2.1 Creating NGINX

Provided:

- nginx.yml file

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  type: ClusterIP
  ports:
    - port: 80
      name: web
  selector:
    app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: ["ReadWriteOnce"]
        resources:
          requests:
            storage: 1Gi
        storageClassName: "standard"
```

This file creates a service and a statefulset. The statefulset creates two pods with nginx image. The service is used to access the pods. The service is a clusterIP service, so it is only accessible from inside the cluster!

2.1.1 Creating a service and statefulset

```
adam@4n6nk8s:~$ vim nginx.yml
adam@4n6nk8s:~$ kubectl apply -f nginx.yml
service/nginx created
statefulset.apps/web created
adam@4n6nk8s:~$ kubectl get pvc
NAME          STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
www-web-0     Bound    pvc-d9ba09bd-e339-4359-b195-133b12a4a0db  1Gi        RWO             standard       45s
www-web-1     Bound    pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb  1Gi        RWO             standard       37s
adam@4n6nk8s:~$ kubectl get pv
NAME          CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM                    STORAGECLASS   REASON   AGE
pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb  1Gi        RWO             Delete            Bound    default/www-web-1       standard   52s
pvc-d9ba09bd-e339-4359-b195-133b12a4a0db  1Gi        RWO             Delete            Bound    default/www-web-0       standard   60s
adam@4n6nk8s:~$ kubectl get po
NAME    READY   STATUS    RESTARTS   AGE
web-0   1/1     Running   0           67s
web-1   1/1     Running   0           59s
adam@4n6nk8s:~$ kubectl get svc
NAME      TYPE        CLUSTER-IP    EXTERNAL-IP   PORT(S)   AGE
kubernetes ClusterIP   10.96.0.1     <none>        443/TCP   8m48s
nginx     ClusterIP   10.101.118.186 <none>        80/TCP    76s
adam@4n6nk8s:~$
```

A service is created with the name `nginx` and a statefulset is created with the name `web`. The statefulset creates two pods with the name `web-0` and `web-1`.

2.1.2 Adding HTML to pods

2.1.2.1 Curling the service internally

```
adam@4n6nk8s:~$ kubectl exec web-1 -- curl 10.101.118.186:80
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
  0    0    0    0    0    0     0      0  0:00:00  0:00:00  0:00:00  0
<head><title>403 Forbidden</title></head>
<body>
<center><h1>403 Forbidden</h1></center>
<hr><center>nginx/1.23.3</center>
</body>
</html>
100 153 100 153    0    0 149k    0  0:00:01  0:00:01  0:00:00 149k
adam@4n6nk8s:~$ kubectl exec web-0 -- curl 10.101.118.186:80
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
  0    0    0    0    0    0     0      0  0:00:00  0:00:00  0:00:00  0
<html>
<head><title>403 Forbidden</title></head>
<body>
<center><h1>403 Forbidden</h1></center>
<hr><center>nginx/1.23.3</center>
</body>
</html>
100 153 100 153    0    0 21857    0  0:00:01  0:00:01  0:00:00 21857
adam@4n6nk8s:~$
```

We got 403 forbidden when trying to curl the homepage of the service from both pods, this is because we don't have any html file being served by NGINX.

Let's check web-0 pod's `nginx/html` contents:

```
adam@4n6nk8s:~$ kubectl exec -it web-0 -- /bin/bash
root@web-0:/# cd /usr/share/nginx/html/
root@web-0:/usr/share/nginx/html# ls -la
total 8
drwxrwxrwx 2 root root 4096 Mar 27 10:12 .
drwxr-xr-x 3 root root 4096 Mar 23 16:00 ..
root@web-0:/usr/share/nginx/html#
```

We can see that the `nginx/html` directory is empty. Let's add some html files to it.

2.1.2.2 Creating a simple hello world homepage

```
adam@4n6nk8s:~$ kubectl exec -it web-0 -- /bin/bash
root@web-0:/# cd /usr/share/nginx/html/
root@web-0:/usr/share/nginx/html# ls -la
total 8
drwxrwxrwx 2 root root 4096 Mar 27 10:12 .
drwxr-xr-x 3 root root 4096 Mar 23 16:00 ..
root@web-0:/usr/share/nginx/html# cat <<EOF > /usr/share/nginx/html/index.html
<html>
<head><title>Hello World from web-0</title></head>
<body>
<center><h1>Hello World from web-0</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
EOF
root@web-0:/usr/share/nginx/html# exit
exit
adam@4n6nk8s:~$ |
```

Let's repeat the process for web-1 pod as well!

```
adam@4n6nk8s:~$ kubectl exec -it web-1 -- /bin/bash
root@web-1:/# cd /usr/share/nginx/html/
root@web-1:/usr/share/nginx/html# ls -la
total 8
drwxrwxrwx 2 root root 4096 Mar 27 10:12 .
drwxr-xr-x 3 root root 4096 Mar 23 16:00 ..
root@web-1:/usr/share/nginx/html#
root@web-1:/usr/share/nginx/html#
root@web-1:/usr/share/nginx/html#
root@web-1:/usr/share/nginx/html# cat <<EOF > /usr/share/nginx/html/index.html
<html>
<head><title>Hello World from web-1</title></head>
<body>
<center><h1>Hello World from web-1</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
EOF
root@web-1:/usr/share/nginx/html#
root@web-1:/usr/share/nginx/html#
root@web-1:/usr/share/nginx/html# exit
exit
adam@4n6nk8s:~$ |
```

2.1.2.3 Curling the service again

```
adam@4n6nk8s:~$ kubectl exec web-0 -- curl 10.101.118.106:80
% Total % Received % Xferd Average Speed Time Time Current
Dload Upload Total Spent Left Speed
100 165 100 165 0 0 161k 0 --:--:-- --:--:-- --:--:-- 161k
<html>
<head><title>Hello World from web-1</title></head>
<body>
<center><h1>Hello World from web-1</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
adam@4n6nk8s:~$ kubectl exec web-0 -- curl 10.101.118.106:80
% Total % Received % Xferd Average Speed Time Time Current
Dload Upload Total Spent Left Speed
0 0 0 0 0 0 0 0 --:--:-- --:--:-- --:--:-- 0<html>
<head><title>Hello World from web-1</title></head>
<body>
<center><h1>Hello World from web-1</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
100 165 100 165 0 0 161k 0 --:--:-- --:--:-- --:--:-- 161k
adam@4n6nk8s:~$ kubectl exec web-1 -- curl 10.101.118.106:80
% Total % Received % Xferd Average Speed Time Time Current
Dload Upload Total Spent Left Speed
100 165 100 165 0 0 161k 0 --:--:-- --:--:-- --:--:-- 161k
<html>
<head><title>Hello World from web-1</title></head>
<body>
<center><h1>Hello World from web-1</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
adam@4n6nk8s:~$ |
```

We keep seeing requests being served by web-1 pod. This is because Kubernetes tend to route requests

to the same pod as long as it is available, for better performane. This is called `sticky sessions`.

In `web-1`, We can rename the `index.html` file to `index2.html` or something, and curl the service again to see that the request is now being served by `web-0` pod, or vice vera.

```
adam@4n6nk8s:~$ kubectl exec -it web-1 -- /bin/bash
root@web-1:/#
root@web-1:/# mv /usr/share/nginx/html/index.html /usr/share/nginx/html/index.html.bak
root@web-1:/#
root@web-1:/# exit
exit
adam@4n6nk8s:~$ |
```

```
adam@4n6nk8s:~$ kubectl exec web-1 -- curl 10.101.118.186:80
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
1<html> 0 0 0 0 0 0 0 0 0 0 0 0
<head><title>Hello World from web-0</title></head>
<body>
<center><h1>Hello World from web-0</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
00 165 100 165 0 0 82500 0 0 0 0 0 82500
adam@4n6nk8s:~$ kubectl exec web-0 -- curl 10.101.118.186:80
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 165 100 165 0 0 41250 0 0 0 0 0 41250
<html>
<head><title>Hello World from web-0</title></head>
<body>
<center><h1>Hello World from web-0</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
adam@4n6nk8s:~$ |
```

As expected, the request is now being served by `web-0` pod.

Let's undo this change for the rest of the lab.

```
adam@4n6nk8s:~$ kubectl exec -it web-1 -- /bin/bash
root@web-1:/#
root@web-1:/# mv /usr/share/nginx/html/index.html.bak /usr/share/nginx/html/index.html
root@web-1:/#
root@web-1:/# exit
exit
adam@4n6nk8s:~$ |
```

2.2 Persistence and Disaster Recovery

In this part, we will be simulating a scheduled outage. We will be removing the statefulset! If we spin up the nginx service and statefulset again, the data will be persisted in the volume claim template, so the new statefulset will have the same data as the old one!

Let's try!

```
adam@4n6nk8s:~$ kubectl delete -f nginx.yml
service "nginx" deleted
statefulset.apps "web" deleted
adam@4n6nk8s:~$ kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
www-web-0 Bound pvc-d9ba09bd-e339-4359-b195-133b12a4a0db 1Gi RWO standard 36m
www-web-1 Bound pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb 1Gi RWO standard 36m
adam@4n6nk8s:~$ kubectl get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS REASON AGE
pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb 1Gi RWO Delete Bound default/www-web-1 standard 36m
pvc-d9ba09bd-e339-4359-b195-133b12a4a0db 1Gi RWO Delete Bound default/www-web-0 standard 36m
adam@4n6nk8s:~$ kubectl get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 44m
adam@4n6nk8s:~$ kubectl get po
No resources found in default namespace.
adam@4n6nk8s:~$
```

Let's recreate the statefulset and service once again...

```
adam@4n6nk8s:~$ kubectl apply -f nginx.yml
service/nginx created
statefulset.apps/web created
adam@4n6nk8s:~$ kubectl get pvc
NAME          STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
www-web-0     Bound     pvc-d9ba09bd-e339-4359-b195-133b12a4a0db 1Gi         RWO             standard       38m
www-web-1     Bound     pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb 1Gi         RWO             standard       38m
adam@4n6nk8s:~$ kubectl get pv
NAME          CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM                STORAGECLASS   REASON   AGE
pvc-2e14dc43-03a9-4a19-847a-4e8c74ae53eb 1Gi         RWO             Delete            Bound    default/www-web-1    standard   38m
pvc-d9ba09bd-e339-4359-b195-133b12a4a0db 1Gi         RWO             Delete            Bound    default/www-web-0    standard   38m
adam@4n6nk8s:~$ kubectl get po
NAME    READY   STATUS    RESTARTS   AGE
web-0   1/1     Running   0           41s
web-1   1/1     Running   0           40s
adam@4n6nk8s:~$ kubectl get svc
NAME    TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes  ClusterIP   10.96.0.1     <none>        443/TCP   46m
nginx     ClusterIP   10.105.238.59 <none>        80/TCP    44s
adam@4n6nk8s:~$
```

And the same `index.html` file is still there!

```
adam@4n6nk8s:~$ kubectl exec -it web-0 -- /bin/bash
root@web-0:/# cat /usr/share/nginx/html/index.html
<html>
<head><title>Hello World from web-0</title></head>
<body>
<center><h1>Hello World from web-0</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
root@web-0:/# exit
exit
adam@4n6nk8s:~$
```

Let's try deleting a pod and see if the data is still there.

```
adam@4n6nk8s:~$ kubectl delete pod web-1 && kubectl get po
pod "web-1" deleted
NAME    READY   STATUS             RESTARTS   AGE
web-0   1/1     Running            0           3m38s
web-1   0/1     ContainerCreating  0           0s
adam@4n6nk8s:~$ kubectl get po
NAME    READY   STATUS    RESTARTS   AGE
web-0   1/1     Running   0           3m46s
web-1   1/1     Running   0           8s
adam@4n6nk8s:~$ kubectl exec -it web-1 -- /bin/bash
root@web-1:/# cat /usr/share/nginx/html/index.html
<html>
<head><title>Hello World from web-1</title></head>
<body>
<center><h1>Hello World from web-1</h1></center>
<hr><center>nginx/1.15.12</center>
</body>
</html>
root@web-1:/#
root@web-1:/# exit
exit
adam@4n6nk8s:~$
```

Indeed, the data is still there! The newly started pod still mount to the same persistent volume as if the pod was never killed.

3 Conclusion

In this lab we learned how to use Minikube to create a local Kubernetes cluster. We also learned how to create a service and a statefulset. We also learned how to persist data in a statefulset using a volume claim template.